

Linear separation and feature maps

DATA 607 — Session 5 — 11/03/2019

Lines in \mathbb{R}^2 : Direction vectors

A line in \mathbb{R}^2 is a set of points of the form

$$L_{\mathbf{u}, \mathbf{v}} := \{\mathbf{u} + t\mathbf{v} : t \in \mathbb{R}\},$$

where

$$\mathbf{u} = \begin{bmatrix} u_1 \\ u_2 \end{bmatrix}, \quad \mathbf{v} = \begin{bmatrix} v_1 \\ v_2 \end{bmatrix} \in \mathbb{R}^2, \quad \mathbf{v} \neq \mathbf{0} = \begin{bmatrix} 0 \\ 0 \end{bmatrix}.$$

Note that $\mathbf{u} = \mathbf{u} + 0\mathbf{v} \in L_{\mathbf{u}, \mathbf{v}}$.

$L_{\mathbf{u}, \mathbf{v}}$ is called the line through \mathbf{u} with direction vector \mathbf{v} .

- $\mathbf{u}' \in L_{\mathbf{u}, \mathbf{v}} \implies L_{\mathbf{u}', \mathbf{v}} = L_{\mathbf{u}, \mathbf{v}}$
- $\mathbf{v}' = c\mathbf{v}, c \in \mathbb{R}, c \neq 0 \implies L_{\mathbf{u}, \mathbf{v}'} = L_{\mathbf{u}, \mathbf{v}}$

Lines in \mathbb{R}^2 : Half-planes; sides; normal vectors

Dot product:

$$\mathbf{u} \cdot \mathbf{v} = \begin{bmatrix} u_1 \\ u_2 \end{bmatrix} \cdot \begin{bmatrix} v_1 \\ v_2 \end{bmatrix} = u_1 v_1 + u_2 v_2 \in \mathbb{R}$$

\mathbf{u} and \mathbf{v} are orthogonal or perpendicular.

\mathbf{w} is a normal vector to $L_{\mathbf{u},\mathbf{v}}$ if $\mathbf{v} \cdot \mathbf{w} = 0$.

Lines in \mathbb{R}^2 : Half-planes and sides

If you delete a line, L , from \mathbb{R}^2 , you're left with two half-planes called the sides of L .

If \mathbf{w} is a nonzero normal vector to $L_{\mathbf{u},\mathbf{v}}$, then the sets

$$H_{\mathbf{u},\mathbf{w}}^- = \{\mathbf{x} : \mathbf{w} \cdot (\mathbf{x} - \mathbf{u}) < 0\} \quad \text{and} \quad H_{\mathbf{u},\mathbf{w}}^+ = \{\mathbf{x} : \mathbf{w} \cdot (\mathbf{x} - \mathbf{u}) > 0\}$$

are the sides of $L_{\mathbf{u},\mathbf{v}}$.

The normal vector \mathbf{w} , plotted with its tail on $L_{\mathbf{u},\mathbf{v}}$, points into $H_{\mathbf{u},\mathbf{w}}^+$.

Linear separation

Let

$$D = \{(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_n, y_n)\}$$

be a dataset, where $\mathbf{x}_i \in \mathbb{R}^2$ and $y_i \in \{-1, 1\}$.

Let

$$D^- = \{(\mathbf{x}_i, y_i) \in D : y_i = -1\}, \quad D^+ = \{(\mathbf{x}_i, y_i) \in D : y_i = +1\}$$

Let L be a line in \mathbb{R}^2 . We say that L separates D if D^- and D^+ are contained in opposite sides of L .

We say that D is linearly separable if there is a line L that separates D .

Not all datasets are linearly separable: The dataset

$$D := \left\{ ((0, 0), 0), ((1, 0), 1), ((1, 1), 0), ((0, 1), 1) \right\}$$

is not linearly separable.

Linear separators need not be unique: The linear separators of

$$D := \left\{ ((0, -1), 0), ((0, 1), 1) \right\}$$

are precisely the lines

$$y = mx + b, \quad b \in (-1, 1).$$

Finding linear separators

Problem: Given a dataset, D , find a vector \mathbf{u} and a nonzero vector \mathbf{w} such that

$$D \cap H_{\mathbf{u}, \mathbf{w}}^+ = D^+ \quad \text{and} \quad D \cap H_{\mathbf{u}, \mathbf{w}}^- = D^-$$

or show that no such \mathbf{u} and \mathbf{w} exist.

We'll begin by analyzing a special case:

Special case: Given a dataset D , find a nonzero vector \mathbf{w} such that

$$D \cap H_{\mathbf{0}, \mathbf{w}}^+ = D^+ \quad \text{and} \quad D \cap H_{\mathbf{0}, \mathbf{w}}^- = D^-$$

$L_{\mathbf{0}, \mathbf{w}}$ does **not** separate D if and only if

$$D^- \cap H_{\mathbf{0}, \mathbf{w}}^+ \neq \emptyset \quad \text{or} \quad D^+ \cap H_{\mathbf{0}, \mathbf{w}}^- \neq \emptyset,$$

...

...or, equivalently, if and only if

$$\left(y_i = -1 \quad \text{and} \quad \mathbf{w} \cdot \mathbf{x}_i > 0 \right) \quad \text{or} \quad \left(y_i = +1 \quad \text{and} \quad \mathbf{w} \cdot \mathbf{x}_i < 0 \right)$$

for some i , or, equivalently, if and only if

$$y_i(\mathbf{w} \cdot \mathbf{x}_i) < 0$$

for some i , or, equivalently, if and only if

$$\min(y_i(\mathbf{w} \cdot \mathbf{x}_i), 0) < 0$$

for some i , or, equivalently,

$$\sum_i \min(y_i(\mathbf{w} \cdot \mathbf{x}_i), 0) < 0,$$

or, equivalently, if and only if

$$\sum_i \max(-y_i(\mathbf{w} \cdot \mathbf{x}_i), 0) > 0.$$

View the term

$$L(\mathbf{w}, \mathbf{x}_i, y_i) := \max(-y_i(\mathbf{w} \cdot \mathbf{x}_i), 0)$$

as a **penalty** or **loss** for \mathbf{x}_i being misclassified by \mathbf{w} , i.e., lying on the wrong side of the line through $\mathbf{0}$ normal to \mathbf{w} .

Assume \mathbf{x}_i is correctly classified, then

$$\left(y_i = -1 \quad \text{and} \quad \mathbf{w} \cdot \mathbf{x}_i < 0 \right) \quad \text{or} \quad \left(y_i = +1 \quad \text{and} \quad \mathbf{w} \cdot \mathbf{x}_i > 0 \right),$$

in which case $y_i(\mathbf{w} \cdot \mathbf{x}_i) > 0$ and

$$-y_i(\mathbf{w} \cdot \mathbf{x}_i) < 0.$$

Thus, the penalty assessed for \mathbf{x}_i being misclassified by \mathbf{w} is

$$L(\mathbf{w}, \mathbf{x}_i, y_i) = \max(-y_i(\mathbf{w} \cdot \mathbf{x}_i), 0) = 0,$$

appropriate since, by hypothesis, \mathbf{x}_i is classified correctly!

Conversely, Assume \mathbf{x}_i is correctly classified, then

$$\left(y_i = -1 \quad \text{and} \quad \mathbf{w} \cdot \mathbf{x}_i > 0\right) \quad \text{or} \quad \left(y_i = +1 \quad \text{and} \quad \mathbf{w} \cdot \mathbf{x}_i < 0\right),$$

in which case $y_i(\mathbf{w} \cdot \mathbf{x}_i) < 0$ and

$$-y_i(\mathbf{w} \cdot \mathbf{x}_i) > 0.$$

Thus, the penalty assessed for \mathbf{x}_i being misclassified by \mathbf{w} is strictly positive:

$$L(\mathbf{w}, \mathbf{x}_i, y_i) = \max(-y_i(\mathbf{w} \cdot \mathbf{x}_i), 0) > -y_i(\mathbf{w} \cdot \mathbf{x}_i).$$

Define the **cost** associated with \mathbf{w} by

$$C(D, \mathbf{w}) = \sum_i L(\mathbf{w}, \mathbf{x}_i, y_i) = \sum_i \max(-y_i(\mathbf{w} \cdot \mathbf{x}_i), 0).$$

Then $L(\mathbf{w})$ separates D if and only if

$$C(D, \mathbf{w}) = 0.$$

$$L(\mathbf{u}, \mathbf{w}, \mathbf{x}_i, y_i) = \sum_i \max(-y_i(\mathbf{w} \cdot (\mathbf{x}_i - \mathbf{u})), 0)$$

$$\begin{aligned} C(D, \mathbf{u}, \mathbf{w}) &= \sum_i L(\mathbf{u}, \mathbf{w}, \mathbf{x}_i, y_i) \\ &= \sum_i \max(-y_i(\mathbf{w} \cdot (\mathbf{x}_i - \mathbf{u})), 0) \end{aligned}$$

Find

$$\operatorname{argmin}_{\mathbf{u}, \mathbf{w}} C(D, \mathbf{u}, \mathbf{w})$$