

Linear Smoothers I

First examples and basic properties

DATA 607 — Session 1 — 25/02/2019

- ① Welcome and introductions
- ② Smoothing
- ③ CODING ACTIVITY 1
 - Code a simple smoother.
 - Explore `sklearn`'s Estimator and Regressor APIs.
- ④ Linear smoothers: First examples
 - k -nearest neighbors
 - local averaging
- ⑤ Evaluating smoothers: Loss and risk
- ⑥ CODING ACTIVITY 2
 - Work with `sklearn`'s built in regressors:
 - `KNeighborsRegressor`, `RadiusNeighborsRegressor`
 - Tune smoothing parameters to minimize average risk.

- 7 Bias-variance decomposition
- 8 CODING ACTIVITY 3
 - Verify the bias-variance decomposition.
- 9 Training error
- 10 Leave one out cross validation
- 11 CODING ACTIVITY 4
 - Tune smoothing parameter using leave one out cross validation.
 - `LeaveOneOut` and `split`
- 12 K -fold cross validation
- 13 CODING ACTIVITY 5
 - Tune smoothing parameter using K -fold cross validation.
 - `cross_val_score`
 - Verify results of CODING ACTIVITY 4

LINEAR SMOOTHERS

Dataset:

$$(\vec{x}_1, Y_1), \dots, (\vec{x}_n, Y_n)$$

Suppose:

$$Y_i = r(\vec{x}_i) + \epsilon_i$$

Definition

A *linear smoother* is an estimator of r of the form

$$\hat{r}(\vec{x}) = \sum_{i=1}^n w_i(\vec{x}) Y_i = \vec{w}(\vec{x}) \cdot \vec{Y}$$

The $w_i(\vec{x})$ are called *weights*.

EXAMPLE 0: LINEAR REGRESSION

Dataset:

$$(\vec{x}_1, Y_1), \dots, (\vec{x}_n, Y_n) \in \mathbb{R}^p \times \mathbb{R}$$

Regression line:

$$\begin{aligned}\hat{r}(\vec{x}) &= \vec{x} \cdot \vec{\beta} \\ &= \vec{x} \cdot ((X^T X)^{-1} X^T \vec{Y}) \\ &= (X(X^T X)^{-1} \vec{x}) \cdot \vec{Y}\end{aligned}$$

This is a linear smoother with $\vec{w}(\vec{x}) = X(X^T X)^{-1} \vec{x}$.

CODING ACTIVITY 1

- Write a nearest neighbor smoother, first as a function and then as a class implementing `sklearn`'s Estimator and Regressor APIs.

EXAMPLE 1: k -NEAREST NEIGHBORS

$N_k(x) :=$ set of the k elements of $\{x_1, \dots, x_n\}$ closest to x .

Definition

The k -nearest neighbor smoother is


$$\hat{r}_k(\vec{x}) = \frac{1}{k} \sum_{i=1}^n \mathbf{1}_{N_k(\vec{x})}(\vec{x}_i) Y_i$$

This is a linear smoother with $w_i(\vec{x}) = \frac{1}{k} \mathbf{1}_{N_k(\vec{x})}(\vec{x}_i)$.


Indicator function of $A \subseteq B$:

$$\mathbf{1}_A(b) := \begin{cases} 1 & \text{if } b \in A, \\ 0 & \text{otherwise.} \end{cases}$$


$k = 2$




$k = 4$




$k = 6$




$k = 8$




$k = 10$




$k = 12$



$k = 14$



$k = 16$



EXAMPLE 2: LOCAL AVERAGING


Definition

The *local average smoother* with *bandwidth* $h > 0$ is


$$\hat{r}_h(\vec{x}) = \frac{\sum_{\|\vec{x}_i - \vec{x}\| < h} Y_i}{\sum_{\|\vec{x}_i - \vec{x}\| < h} 1}$$

In words: $r_h(\vec{x})$ is the average of the Y_i for which \vec{x}_i lies a distance less than h from \vec{x} .


$h = 0.025$




$h = 0.05$




$h = 0.075$




$h = 0.1$




$h = 0.125$




$h = 0.15$



$h = 0.175$



$h = 0.2$



OVERFITTING AND UNDERFITTING

For small k or h , the graphs of $\hat{r}(x)$ are too wiggly. We've *overfit* or *undersmoothed* the data.

For big k or h , the graphs of $\hat{r}(x)$ data very well (they're too flat). We say we've *underfit* or *oversmoothed* the data.

How do we choose k or h optimally?

EVALUATING SMOOTHERS: LOSS AND RISK

Definitions

- ① *Squared error loss at x :*

$$L(\hat{r}(x), r(x)) := (\hat{r}(x) - r(x))^2$$

- $L(\hat{r}(x), r(x))$ is a random variable as $\hat{r}(x)$ is.

- ② *Mean squared error (MSE) or risk at x :*

$$R(\hat{r}(x), r(x)) := \mathbb{E}[L(\hat{r}(x), r(x))] \approx \frac{1}{n} \sum_{i=1}^n L(\hat{r}(x_i), r(x_i))$$

- ③ *Average MSE or average risk:*

$$R(\hat{r}, r) := \int R(\hat{r}(x), r(x)) dx \approx \frac{1}{n} \sum_{i=1}^n R(\hat{r}(x_i), r(x_i))$$

CHOOSING THE SMOOTHING PARAMETER

Obviously, smoothers with smaller average risk are preferred.

So, choose the smoothing parameter to *minimize average risk*:

$$k := \operatorname{argmin}_k R(\hat{r}_k, r), \quad h := \operatorname{argmin}_h R(\hat{r}_h, r)$$

But there's a problem...

WE DON'T KNOW $r(x)$!

Loss at x :
$$L(\hat{r}(x), r(x)) = (\hat{r}(x) - r(x))^2$$

Risk at x :
$$R(\hat{r}(x), r(x)) = \mathbb{E}[L(\hat{r}(x), r(x))]$$

Average risk at x :
$$R(\hat{r}, r) = \int R(\hat{r}(x), r(x)) dx$$

These expressions all involve $r(x)$, which we typically don't know!

We can't compute these expressions; they also need to be estimated.

CODING ACTIVITY 2

But, we can work in an artificial situation where we *do* know $r(x)$:

- Choose:
 - a function, $r(x)$
 - x_1, \dots, x_n
 - a fine partition, P , of an interval containing the x_i
 - k (or h)
- for $k = 1, \dots, K$
 - for $j = 0, \dots, J - 1$:
 - Generate $y_i^{(j)}$, $1 \leq i \leq n$, by sampling from $N(r_k(x_i), \sigma^2)$.
 - Compute the losses $L(\hat{r}_k^{(j)}(x), r(x))$ at each x in P .
 - Average the losses over j to get the risks, $R(\hat{r}_k(x), r(x))$.
 - Average the risks over x to get the average risk, $R(\hat{r}_k, r)$.
- Plot $R(\hat{r}_k, r)$ vs k and choose k .

THE BIAS-VARIANCE DECOMPOSITION

Definition

The *bias* of $\hat{r}(x)$, as an estimator of $r(x)$, is

$$\text{Bias}(\hat{r}(x), r(x)) := \mathbb{E}[\hat{r}(x)] - r(x)$$


Bias-variance decomposition

$$R(\hat{r}(x), r(x)) = \text{Var } \hat{r}(x) + \text{Bias}(\hat{r}(x), r(x))^2$$


$\text{Var } \hat{r}(x)$ big / graph of \hat{r} wiggly / overfit / undersmoothed

$\text{Bias}(\hat{r}(x), r(x))$ big / graph of \hat{r} too flat / underfit / oversmoothed

$x = 0.25$



$x = 0.75$



Variance decreases as k increases.

Bias increases as k increases.

Define:

$$\text{Var } \hat{r} := \int R(\hat{r}(x), r(x)) dx, \quad \text{Bias}(\hat{r}, r)^2 := \int \text{Bias}(\hat{r}(x), r(x))^2 dx$$

Integrate the bias-variance decomposition:

$$R(\hat{r}, r) = \text{Var } \hat{r} + \text{Bias}(\hat{r}, r)^2$$

Since r is unknown, these quantities can only be estimated.

CODING ACTIVITY 3

- For our running synthetic example, plot

$$R(\hat{r}_k, r), \text{Var } \hat{r}_k, \text{ and } \text{Bias}(\hat{r}_k, r)^2$$

versus k on the same axes.

- Verify the integrated bias-variance decomposition using your computed $R(\hat{r}_k, r)$, $\text{Var } \hat{r}_k$, and $\text{Bias}(\hat{r}_k, r)^2$.

Reuse as much of your code from CODING ACTIVITY 2 as possible.

Definition

The *empirical risk* or *training error* is

$$\frac{1}{n} \sum_{i=1}^n L(\hat{r}(x_i), Y_i).$$

Empirical risk is not good estimator of R : We trained \hat{r} on the (x_i, Y_i) , making $L(\hat{r}(x_i), Y_i)$ biased downwards.

LEAVE ONE OUT CROSS-VALIDATION

Definition

The *leave one out cross validation (LOOCV)* score is

$$\hat{R}(h) := \frac{1}{n} \sum_{i=1}^n (Y_i - \hat{r}_h^{(-i)}(x_i))^2,$$

where $\hat{r}_h^{(-i)}$ is computed using the subdataset of the original one by removing the point (x_i, Y_i) .

$\hat{R}(h)$ is typically a good estimator of the average risk of \hat{r}_h :

$$\hat{R}(h) \approx R(\hat{r}_h, r) + \sigma^2, \quad \text{where } \sigma^2 = \text{Var } Y_i.$$

We choose our smoothing parameter to be the one that minimizes $\hat{R}(h)$:

$$h := \operatorname{argmin}_h \hat{R}(h)$$

CODING ACTIVITY 4

For our running example, plot $\hat{R}(k)$ versus k . Compute the $\hat{R}(k)$ using `sklearn`'s `LeaveOneOut` class and the generator returned by its `split` method.

Which k that minimizes $\hat{R}(k)$? Compare with the results of CODING ACTIVITY 2.

For this optimal value of k , plot $\hat{r}_k^{(-i)}(x_i)$ versus x_i and $r(x)$ versus x on the same axes.

Remark: `LeaveOneOut` is fairly low-level; `sklearn` provides more convenient ways to tune hyperparameters using cross validation. Sometimes, though, it's necessary to work with the lower-level constructs.

K-FOLD CROSS VALIDATION

Partition $\{1, \dots, n\}$ into K folds, I_1, \dots, I_K , of roughly equal size.

Let $\hat{r}_h^{(-I_j)}$ be the smoother computed using the subdataset of the original one obtained by removing (x_i, Y_i) , for $i \in I_j$.

Definition

The K -fold cross validation score is

$$\hat{R}_K(h) := \sum_{j=1}^K \frac{1}{|I_j|} \sum_{i \in I_j} (Y_i - r_h^{(-I_j)}(x_i))^2$$

We can use \hat{R}_K in place of \hat{R} to select a smoothing parameter.

In practice, K is usually 5 (sklearn's default) or 10.

CODING ACTIVITY 5

Find the values of k that minimize the 3-, 5-, and 10-fold cross validation scores. Compute these scores using `cross_val_score` from `sklearn.model_selection`.

Confirm your results from CODING ACTIVITY 4 by performing LOOCV as n -fold cross validation, n being the size of the dataset.

BINARY CLASSIFICATION

Regression:

$$Y_i = r(x_i) + \epsilon_i$$

Binary classification:

$$Y_i = r(x_i) \in \{0, 1\}$$

Want \hat{r} to be $\{0, 1\}$ -valued, too.

If $w_i(x) \geq 0$ and $\sum_i w_i(x) = 1$ for all x ,

$$\hat{r}(x) = d\left(\sum_{i=1}^n w_i(x) Y_i\right) \in \{0, 1\}$$

for any *decision function* $d : [0, 1] \rightarrow \{0, 1\}$.

Most common: $d(x) = \mathbf{1}_{[1/2, 1]}(x)$ or $d(x) = \mathbf{1}_{[\alpha, 1]}(x)$.

$$L(\hat{r}(x), r(x)) = |\hat{r}(x) - r(x)| = \begin{cases} 1 & \text{if } \hat{r}(x) = r(x), \\ 0 & \text{if } \hat{r}(x) \neq r(x) \end{cases}$$

$$R(\hat{r}(x), r(x)) = \mathbb{E}[L(\hat{r}(x), r(x))]$$

$$R(\hat{r}, r) = \int R(\hat{r}(x), r(x)) dx$$